


Backpropagation in Deep Learning

- Backpropagation (short for "backward propagation of errors") is a key algorithm used for training artificial neural networks.
- It's a supervised learning technique that helps the network learn by adjusting the weights of neurons in order to minimize the difference between predicted outputs and actual target values.


📌 Key Idea:

1. **Forward Propagation** → Compute predictions.
2. **Calculate Loss** → **Calculating the error** between predictions and true values.
3. **Backward Propagation** → Compute **gradients (derivatives)** of the loss function w.r.t. weights.
4. **Update Weights** → Adjust weights using Gradient Descent to reduce the error.





Backpropagation finds out appropriate values of weights & biases.



Data:

iq	cgpa	lpa
80	8	3
60	9	5
70	5	8
120	7	11

Steps

Note: Here, activation function is **Linear**.

1. Initialize weights & biases (w, b)

$w \rightarrow 1$

$b \rightarrow 0$

2. Select a row

- Feed the data
 - eg. **80 & 8** (Row 1 🙌)

3. Predict using the initial weight & bias

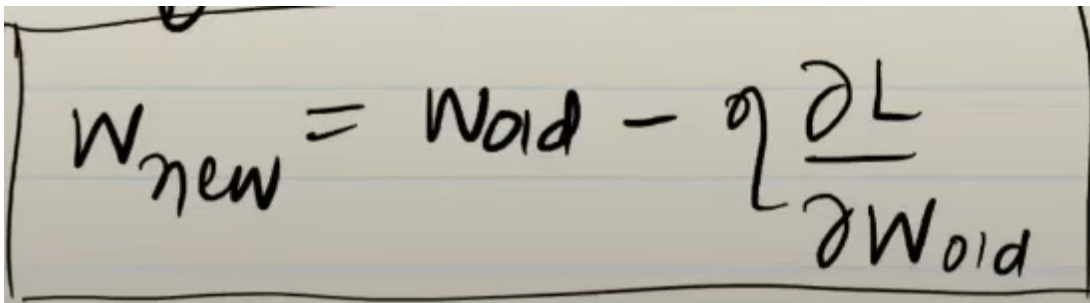
- This step is called **forward propagation**.
- eg. The model predicted **18**
 - 🙌 This is obviously far from the actual value (**3**)

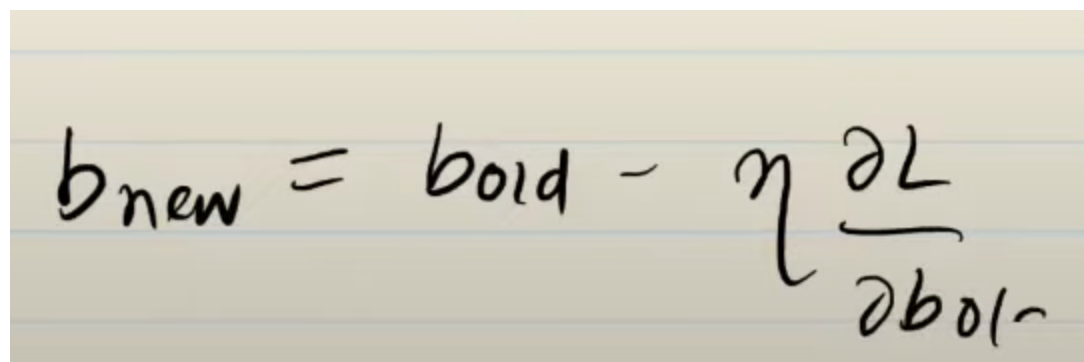
4. Choose a **Loss Function**

- eg. **MSE** for regression
 - $(3 - 18)^2 = 225$
- Here, we have to decrease the value.
- To decrease the value, we have to **go back and change the values** of w & b
 - Therefore, the name → **Back propagation**

5. Update the weights & biases

- Use → **Gradient Descent**


$$w_{\text{new}} = w_{\text{old}} - \eta \frac{\partial L}{\partial w_{\text{old}}}$$


$$b_{\text{new}} = b_{\text{old}} - \eta \frac{\partial L}{\partial b_{\text{old}}}$$

η = Learning rate

For above example, we have to calculate these 9 derivatives:

$$\frac{\partial L}{\partial w_{11}^2}, \frac{\partial L}{\partial w_{21}^2}, \frac{\partial L}{\partial b_{21}} \quad \left| \quad \frac{\partial L}{\partial w_{11}^1}, \frac{\partial L}{\partial w_{21}^1}, \frac{\partial L}{\partial b_{11}} \quad \left| \quad \frac{\partial L}{\partial w_{12}^1}, \frac{\partial L}{\partial w_{22}^1}, \frac{\partial L}{\partial b_{12}} \right.$$

Meaning of Derivative:



How much the **L** changes when we are changing *w* or *b*

- L is loss
- The loss changes when we make changes in weights and biases. We're finding out → how much does it change?

We cannot directly calculate:

$$\frac{\partial L}{\partial \hat{y}}$$

First we have to calculate

$$\frac{\partial \hat{y}}{\partial w_{11}^2}$$

- After that we'll multiply these 2

$$\frac{\partial L}{\partial w_{11}^2} = \left[\frac{\partial L}{\partial \hat{y}} \right] \times \left[\frac{\partial \hat{y}}{\partial w_{11}^2} \right]$$

👉 This is called chain rule.

$$\left[\frac{\partial L}{\partial \hat{y}} \right] = \frac{\partial}{\partial \hat{y}} (y - \hat{y})^2 = [-2(y - \hat{y})]$$

$$\frac{\partial \hat{y}}{\partial w_{11}^2} = \frac{\partial}{\partial w_{11}^2} [0_{11} w_{11}^2 + 0_{12} w_{21}^2 + b_{21}]$$

$$= 0_{11}$$

$$\frac{\partial L}{\partial w_{11}^2} = -2(y - \hat{y})o_{11}$$

Similarly, the second derivative will be

$$\frac{\partial L}{\partial w_{21}^2} = -2(y - \hat{y})o_{12}$$

Bias:

$$\frac{\partial L}{\partial b_{21}} = -2(y - \hat{y})$$

6. Repeat this process

💡 Key Insights

1. **Vanishing Gradients:**

- If gradients become too small (e.g., in sigmoid/tanh), early layers learn **slowly**.
- **Fix:** Use **ReLU** or **batch normalization**.

2. Exploding Gradients:

- If gradients grow too large (common in RNNs), training becomes unstable.
- **Fix:** Use **gradient clipping**.

3. Local Minima:

- Backpropagation can get stuck in suboptimal solutions.
- **Fix:** Use **momentum** (e.g., Adam optimizer).

MLP Memoization – Optimizing Neural Network Efficiency



Memoization is already built-in in Keras library

- **Memoization** (caching intermediate results) can **speed up training & inference** in **Multilayer Perceptrons (MLPs)** by avoiding redundant computations.



What is Memoization in MLPs?

- **Stores layer outputs** during forward/backward passes.
- **Reuses them** instead of recalculating (e.g., in loops or repeated calls).
- **Trade-off:** Saves compute but increases memory usage.

Limitations:

- **Memory usage:** Memoization consumes memory to store intermediate results, which could become an issue when training deep networks with large datasets.
- **Batch processing:** When processing large batches of data, memoization might be less effective because the inputs are often different across samples, leading to fewer cache hits.