

# Early Stopping & Data Scaling

## Early Stopping

```
from keras.callbacks import EarlyStopping  
  
model.fit(callbacks=[early_stopping])
```

- **Early Stopping** is a regularization technique that halts training when the model stops improving on a **validation set**, **preventing overfitting and saving computation time**.

### How Early Stopping Works

1. **Monitor Validation Metric** (e.g., `val_loss`, `val_accuracy` ).
2. **Stop Training** if the metric doesn't improve for `N` epochs ( `patience` ).
3. **Restore Best Weights** (optional) to the epoch with the optimal performance.

### Key Parameters:

Parameter	Description	Recommended Value
<code>monitor</code>	Metric to track (e.g., <code>val_loss</code> , <code>val_accuracy</code> )	<code>val_loss</code> (default)
<code>patience</code>	Epochs to wait before stopping	5-10
<code>restore_best_weights</code>	Revert to best model weights	True
<code>min_delta</code>	Minimum change to qualify as improvement	0.001
<code>mode</code>	<code>auto</code> / <code>min</code> / <code>max</code> (e.g., minimize loss or maximize accuracy)	<code>auto</code>

## Implementation in Keras

```

from keras.callbacks import EarlyStopping

early_stopping = EarlyStopping(
    monitor='val_loss', # Track validation loss
    patience=5,         # Stop after 5 epochs without improvement
    restore_best_weights=True, # Revert to best weights
    min_delta=0.001     # Minimum change to count as improvement
)

model.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    epochs=100,
    callbacks=[early_stopping] # Add to training
)

```

- **monitor** : What metric to monitor ( 'val\_loss' , 'val\_accuracy' , etc.).
- **patience** : The number of epochs to wait for improvement before stopping.
- **min\_delta** : Minimum change in the monitored metric to count as an improvement.
- **restore\_best\_weights** : Ensures the model returns to the weights from the epoch with the best performance on the validation set.

## When to Use Early Stopping

Early stopping is especially useful when:

- You don't have a large amount of training data.
- You're working with deep learning models that have a risk of overfitting due to their complexity.
- You want to ensure that your model generalizes well without manually tuning for too many epochs.


## Why Use Early Stopping?

- ✓ **Prevents Overfitting:** Stops training before the model memorizes noise.
- ✓ **Saves Time:** Avoids unnecessary epochs.
- ✓ **No Extra Parameters:** Unlike L2/Dropout, it doesn't modify the model.

## Data Scaling (Normalization)

Neural networks are sensitive to input scale. Proper data scaling ensures:

- ✓ Faster convergence (optimizers work efficiently).
- ✓ Better performance (avoids dominance of large-scale features).
- ✓ Numerical stability (prevents exploding/vanishing gradients).



## Standardize vs Normalize

Handwritten notes on lined paper showing two scaling methods:

- Standard:**  $\frac{X_i - \mu}{\sigma}$
- Normalize:**  $\frac{X_i - X_{min}}{X_{max} - X_{min}}$

- **Standardize** → -1 to 1
- **Normalize** → 0 to 1

### When to standardise & when to normalize?

- If you know the max & min values → Normalize
- Otherwise → Standardise
  - Standardise if data is normally distributed.

### Standardization (Z-Score Normalization)

$$X_{\text{scaled}} = \frac{X - \mu}{\sigma}$$

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## Min-Max Normalization

$$X_{\text{scaled}} = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## When to Scale Data?

Algorithm	Scaling Needed?	Recommended Method
Neural Networks	✓ Yes	Standardization
CNNs (Images)	✓ Yes	Min-Max (0-1)
RNNs (Time Series)	✓ Yes	Standardization
Tree-Based Models	✗ No	Not required

## Common Mistakes

✗ **Scaling Test Data Separately:** Always use `transform()` with the **same scaler** from training.

✗ **Ignoring Categorical Features:** Scale only **numerical** features (one-hot encode categoricals).

**✗ Forgetting to Scale Outputs:** In regression, scale  $yy$  if it has a large range.