


Gradient Descent in Neural Networks

Batch Gradient Descent

- **Batch GD** uses the entire data to calculate slope in this equation.
 - **Fast**

```
model.fit(X_scaled,y,epochs=500,batch_size=n,validation_split=0.2)
```

- For batch Batch GD, `batch_size = n` → **No. of rows**
 - Here, you're creating a single batch of n rows
- Loss decreases **smoothly** in Batch GD



Stochastic Gradient Descent

- **Stochastic Gradient Descent** takes 1 row & calculates values for each row

- **Slow**


```
model.fit(X_scaled,y,epochs=500,batch_size=1,validation_split=0.2)
```

- For batch SGD, `batch_size = 1`



Stochastic GD gives **better accuracy** with **same no. of epochs**.

- Loss decreases jaggedly in SGD



- 🖐️ This jaggedness helps SGD to move out of local minima.
- But the solution is also approximate.

Mini-Batch Gradient Descent

- This balances the Batch and Stochastic GD
- You define a batch size - eg. 30

- It gets updated when it goes through 30 rows
- If there are 300 rows, **Mini batch Descent** it will update **10 times**
- **Stochastics GD** would update → **300 times**
- **Batch GD** would update → **1 Time**



MOSTLY USED GD → Mini-Batch GD

Interview Question

Q. Why batch-size is provided in multiples of 2?

- RAM is designed to handle binary values.
- It's done for effective use of RAM.
- Optimization technique

What if batch-size doesn't divide # rows properly

e.g # of rows $n = 400$
batch-size = 150

$$\# \text{ of batch} = \frac{400}{150} = 2.66$$

150, 150, left 100
↑ ↑ ↑
1 batch 2 batch 3rd batch