

Hyperparameter Tuning a Neural Network

Library → Keras tuner

```
pip install keras-tuner
```

```
import keras_tuner as kt
```


Dataset → *Pima Indians Diabetes Database*

```
df = pd.read_csv(r'https://raw.githubusercontent.com/npradaschnor/Pima-Indians-Diabetes-Dataset/refs/heads/master/diabetes.csv')
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

```
sns.heatmap(df.corr(),annot=True)
```



```
df.corr().Outcome
```

```
Pregnancies      0.221898
Glucose           0.466581
BloodPressure    0.065068
SkinThickness    0.074752
Insulin          0.130548
BMI              0.292695
DiabetesPedigreeFunction 0.173844
Age              0.238356
Outcome          1.000000
Name: Outcome, dtype: float64
```

```
X= df.drop(columns=['Outcome'])
y= df['Outcome']
```

```
# SCALE
```

```
from sklearn.preprocessing import StandardScaler as sc
```

```
scaler = sc()
```

```
X= scaler.fit_transform(X)
```

```
# train_test_split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state
=1)
```

Build a model

```
import tensorflow
```

```
from tensorflow import keras
```

```
from keras import Sequential
```

```
from keras.layers import Dense
```

```

model = Sequential()
model.add(Dense(32, activation='relu',input_dim=8))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='BinaryCrossentropy', metrics=['accuracy'])

model.fit(X_train,y_train, batch_size=32, epochs=10,validation_data=(X_test, y_test))

```

```

Epoch 4/10
20/20 ————— 0s 5ms/step - accuracy: 0.7613 - loss: 0.4407 - val_accuracy: 0.8117 - val_loss: 0.4729
Epoch 5/10
20/20 ————— 0s 5ms/step - accuracy: 0.7836 - loss: 0.4345 - val_accuracy: 0.8052 - val_loss: 0.4738
Epoch 6/10
20/20 ————— 0s 6ms/step - accuracy: 0.7972 - loss: 0.4314 - val_accuracy: 0.8052 - val_loss: 0.4743
Epoch 7/10
20/20 ————— 0s 5ms/step - accuracy: 0.7925 - loss: 0.4083 - val_accuracy: 0.7987 - val_loss: 0.4756
Epoch 8/10
20/20 ————— 0s 6ms/step - accuracy: 0.8057 - loss: 0.4033 - val_accuracy: 0.8052 - val_loss: 0.4749
Epoch 9/10
20/20 ————— 0s 6ms/step - accuracy: 0.7557 - loss: 0.4786 - val_accuracy: 0.8182 - val_loss: 0.4710
Epoch 10/10
20/20 ————— 0s 7ms/step - accuracy: 0.7819 - loss: 0.4399 - val_accuracy: 0.8117 - val_loss: 0.4696

```

Tune the Hyperparameters

```
pip install keras-tuner
```

```
import keras_tuner as kt
```

Steps:

1. Make a function
2. Pass the function into a tuner object

Find best optimizer:

```
def build_model(hp):
    model = Sequential()
    model.add(Dense(32, activation='relu',input_dim=8))
    model.add(Dense(1, activation='sigmoid'))

    optimizer = hp.Choice('optimizer', ['adam','sgd','rmsprop','adadelta'])

    model.compile(optimizer= optimizer, loss='binary_crossentropy', metrics=['a
ccuracy'])
    return model
```

Make Tuner Object:

```
tuner = kt.RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=5)
```

Default Objective: `'val_loss'`

- Option 2 : `'val_accuracy'` : The validation accuracy.

`max_trials` : Default = 10

```
tuner.search(X_train,y_train, epochs=5, validation_data=(X_test,y_test))
```

```
Trial 4 Complete [00h 00m 02s]
val_accuracy: 0.7597402334213257

Best val_accuracy So Far: 0.798701286315918
Total elapsed time: 00h 00m 10s
```

Get best HP:

```
tuner.get_best_hyperparameters()[0].values
```

```
{'optimizer': 'rmsprop'}
```

Get best model

```
model = tuner.get_best_models(num_models=1)[0]  
model.summary()
```

```
...  
| Layer (type) | Output Shape | Param # |  
|-----|-----|-----|  
| dense (Dense) | (None, 32) | 288 |  
| dense_1 (Dense) | (None, 1) | 33 |  
...  
Total params: 321 (1.25 KB)  
...  
Trainable params: 321 (1.25 KB)  
...  
Non-trainable params: 0 (0.00 B)
```

Fit :

```
model.fit(X_train,y_train,batch_size=32,epochs=100,initial_epoch=6,validation  
_data=(X_test,y_test))
```

`initial_epoch=6` because we have already run 5 epochs

```

Epoch 10/100 0s 5ms/step - accuracy: 0.7521 - loss: 0.4747 - val_accuracy: 0.7792 - val_loss: 0.4661
Epoch 19/100
...
Epoch 99/100
20/20 0s 5ms/step - accuracy: 0.8175 - loss: 0.3862 - val_accuracy: 0.8052 - val_loss: 0.4631
Epoch 100/100
20/20 0s 5ms/step - accuracy: 0.7963 - loss: 0.4181 - val_accuracy: 0.8117 - val_loss: 0.4630

```

Find out number of neurons:

```

def build_model(hp):
    model= Sequential()
    units = hp.Int('units', 8,128,8)
    model.add(Dense(units=units,activation='relu', input_dim=8))
    model.add(Dense(1,activation='sigmoid'))

    model.compile(optimizer='rmsprop', loss='binary_crossentropy',metrics=['a
ccuracy'], )

    return model

```

8,128,8 → min_value = 8,max_value = 128, step=8

Create Tuner Object:

```

tuner = kt.RandomSearch(build_model, objective='val_accuracy', max_trials=
5, directory='mydir', project_name='kuch_bhi')

```

Search best parameters:

```

tuner.search(X_train,y_train, validation_data=(X_test,y_test))

```

```
Trial 5 Complete [00h 00m 02s]
val_accuracy: 0.7857142686843872

Best val_accuracy So Far: 0.8116883039474487
Total elapsed time: 00h 00m 12s
```

```
tuner.get_best_hyperparameters()[0].values
```

```
{'units': 112}
```

Best model:

```
model2= tuner.get_best_models(num_models=1)[0]
```

Fit again:

```
model2.fit(X_train,y_train,batch_size=32,epochs=100,initial_epoch=6, validation_data=(X_test,y_test))
```

```
20/20 ————— 0s 2ms/step - accuracy: 0.7743 - loss: 0.4417
Epoch 16/100
20/20 ————— 0s 3ms/step - accuracy: 0.7871 - loss: 0.4485
Epoch 17/100
20/20 ————— 0s 2ms/step - accuracy: 0.7838 - loss: 0.4407
Epoch 18/100
20/20 ————— 0s 2ms/step - accuracy: 0.7875 - loss: 0.4444
Epoch 19/100
...
Epoch 99/100
20/20 ————— 0s 2ms/step - accuracy: 0.8247 - loss: 0.3827
Epoch 100/100
20/20 ————— 0s 2ms/step - accuracy: 0.8131 - loss: 0.4134
```

Select No. of layers:


```
def build_model(hp):
    model = Sequential()
    model.add(Dense(112, activation='relu', input_dim=8))

    for i in range(hp.Int('num_layers', min_value=1, max_value=10)):
        model.add(Dense(72, activation='relu'))
        model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])

    return model
```

```
for i in range(hp.Int('num_layers', min_value=1, max_value=10))
```

- There will be 10 models with 1 to 10 layers

`hp.Int('num_layers', min_value=1, max_value=10):` This is where the magic of hyperparameter tuning comes in!

- `hp.Int` **means pick a whole number (integer) for the hyperparameter named 'num_layers'**
- `min_value=1, max_value=10` means the number of layers can be between 1 and 10
- `for i in range(...)` : This loop will run as many times as the number of layers chosen by `hp.Int`. For example, if `hp.Int` picks 3, the loop runs 3 times to add 3 layers.

```
tuner=kt.RandomSearch(build_model, objective='val_accuracy', max_trials=5,
    directory= 'mydir', project_name= 'num_layers')
```

```
tuner.search(X_train,y_train, epochs=5, validation_data=(X_test,y_test))
```

```
Trial 5 Complete [00h 00m 04s]
val_accuracy: 0.798701286315918

Best val_accuracy So Far: 0.8181818127632141
Total elapsed time: 00h 00m 17s
```

```
tuner.get_best_hyperparameters()[0].values
```

```
{'num_layers': 4}
```

```
model3 = tuner.get_best_models(num_models=1)[0]
```

```
model3.fit(X_train,y_train,epochs=100,initial_epoch=6,validation_data=(X_test,
y_test))
```

```
Epoch 16/100
20/20 ————— 0s 5ms/step - accuracy: 0.8639 - loss: 0.3225 - val_accuracy: 0.7727 - val_loss: 0.4954
Epoch 17/100
20/20 ————— 0s 5ms/step - accuracy: 0.8186 - loss: 0.3993 - val_accuracy: 0.7857 - val_loss: 0.4888
Epoch 18/100
20/20 ————— 0s 5ms/step - accuracy: 0.8488 - loss: 0.3683 - val_accuracy: 0.7727 - val_loss: 0.5105
Epoch 19/100
...
Epoch 99/100
20/20 ————— 0s 5ms/step - accuracy: 0.9986 - loss: 0.0056 - val_accuracy: 0.7532 - val_loss: 2.0436
Epoch 100/100
20/20 ————— 0s 5ms/step - accuracy: 0.9933 - loss: 0.0247 - val_accuracy: 0.7273 - val_loss: 2.3232
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

👉 Validation Accuracy is decreasing

Now find of both No. of layers & Nodes

```
def build_model(hp):
```

```

model = Sequential()
counter = 0

for i in range(hp.Int('num_layers',min_value=1,max_value=10)):

    if counter == 0:
        model.add(Dense(
            hp.Int('units'+str(i), min_value=8, max_value=128, step=8),
            activation= hp.Choice('activation' + str(i), values=['relu','tanh','sigmoid']),
            input_dim=8
        ))
    else:
        model.add(Dense(hp.Int('units' + str(i), min_value=8, max_value=128,step=
8),activation= hp.Choice('activation' + str(i), values=['relu','tanh','sigmoid'])))
        counter += 1

    model.add(Dense(1,activation='sigmoid'))

    model.compile(optimizer=hp.Choice('optimizer',values=['rmsprop','adam','s
gd','nadam','adadelta']),
        loss='binary_crossentropy',
        metrics=['accuracy'])

return model

```

Tuner Object:

```

tuner = kt.RandomSearch(build_model,
                        objective='val_accuracy',
                        max_trials=5,
                        directory="mydir",
                        project_name='final')

```

```
tuner.search(X_train,y_train,epochs=5,validation_data=(X_test,y_test))
```

```
Trial 5 Complete [00h 00m 02s]
val_accuracy: 0.6428571343421936

Best val_accuracy So Far: 0.6428571343421936
Total elapsed time: 00h 00m 27s
```

```
tuner.get_best_hyperparameters()[0].values
```

```
{'num_layers': 7,
 'units0': 72,
 'activation0': 'tanh',
 'optimizer': 'nadam',
 'units1': 8,
 'activation1': 'relu',
 'units2': 8,
 'activation2': 'relu',
 'units3': 8,
 'activation3': 'relu',
 'units4': 8,
 'activation4': 'relu',
 'units5': 8,
 'activation5': 'relu',
 'units6': 8,
 'activation6': 'relu'}
```

```
model4= tuner.get_best_models(num_models=1)[0]
model4.fit(X_train,y_train,epochs=200,initial_epoch=6,validation_data=(X_test,
y_test))
```

```

Epoch 16/200 0s 7ms/step - accuracy: 0.6555 - loss: 0.6443 - val_accuracy: 0.6429 - val_loss: 0.6518
Epoch 17/200 0s 7ms/step - accuracy: 0.6522 - loss: 0.6463 - val_accuracy: 0.6429 - val_loss: 0.6518
Epoch 18/200 0s 7ms/step - accuracy: 0.6633 - loss: 0.6395 - val_accuracy: 0.6429 - val_loss: 0.6518
Epoch 19/200 ...
Epoch 199/200 0s 7ms/step - accuracy: 0.7936 - loss: 0.4690 - val_accuracy: 0.7532 - val_loss: 0.5300
Epoch 200/200 0s 6ms/step - accuracy: 0.7482 - loss: 0.5029 - val_accuracy: 0.7532 - val_loss: 0.5292
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

```

Add dropout layer:

- All code is same as above. just added 📌

```

model.add(Dropout(hp.Choice('dropout' + str(i), values=[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9])))

```

```

def build_model(hp):

    model = Sequential()

    counter = 0

    for i in range(hp.Int('num_layers',min_value=1,max_value=10)):

        if counter == 0:

            model.add(Dense(hp.Int('units' + str(i), min_value=8, max_value=128,step=8),activation= hp.Choice('activation' + str(i), values=['relu','tanh','sigmoid']),input_dim=8))

            model.add(Dropout(hp.Choice('dropout' + str(i), values=[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9])))

        else:

            model.add(Dense(hp.Int('units' + str(i), min_value=8, max_value=128,step=8),activation= hp.Choice('activation' + str(i), values=['relu','tanh','sigmoid'])))

```

```

        model.add(Dropout(hp.Choice('dropout' + str(i), values=[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9])))

        counter+=1

    model.add(Dense(1,activation='sigmoid'))

    model.compile(optimizer=hp.Choice('optimizer',values=['rmsprop','adam','sgd','nadam','adadelat']),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

    return model

```

counter = 0

- This sets up a variable called counter and starts it at 0. It's like a little counter to keep track of which layer we're adding (first, second, third, etc.).
- We'll use this to decide if we're adding the first layer (which needs input_dim) or a later layer.

'units' + str(i) :

- This makes a unique name for each layer's hyperparameter (like 'units0', 'units1', etc.), so the tuner can track them separately

Why **str(i)** ?

- It gives each layer a unique name (like 'units0', 'units1') so the tuner can pick different settings for each layer.

Can We Simplify Without **str(i)** ?

- Yes, but then all layers will have the same settings (same neurons, activation, dropout), which is less flexible.
- Simple answer → **NO**.

Create a tuner object:

```
tuner = kt.RandomSearch(build_model,  
                        objective='val_accuracy',  
                        max_trials=3,  
                        directory='mydir',  
                        project_name='kuch bhi')
```

```
tuner.search(X_train,y_train,epochs=5,validation_data=(X_test,y_test))
```

```
Trial 3 Complete [00h 00m 03s]  
val_accuracy: 0.7142857313156128  
  
Best val_accuracy So Far: 0.7142857313156128  
Total elapsed time: 00h 00m 12s
```

```
tuner.get_best_hyperparameters()[0].values
```

```
{'num_layers': 2,  
  'units0': 24,  
  'activation0': 'relu',  
  'dropout0': 0.5,  
  'optimizer': 'sgd',  
  'units1': 32,  
  'activation1': 'relu',  
  'dropout1': 0.6,  
  'units2': 32,  
  'activation2': 'relu',  
  'dropout2': 0.7,  
  'units3': 72,  
  'activation3': 'tanh',  
  'dropout3': 0.5,  
  'units4': 104,  
  'activation4': 'relu',  
  'dropout4': 0.6,  
  'units5': 64,  
  'activation5': 'tanh',  
  'dropout5': 0.7,  
  'units6': 80,  
  'activation6': 'sigmoid',  
  'dropout6': 0.3,  
  'units7': 56,  
  'activation7': 'tanh',  
  'dropout7': 0.9,  
  'units8': 56,  
  'activation8': 'sigmoid',  
  'dropout8': 0.7}
```

```
model = tuner.get_best_models(num_models=1)[0]
```

Fit

```
model.fit(X_train,y_train,epochs=200,initial_epoch=6,validation_data=(X_test,y_test))
```



```
Epoch 16/200
20/20 — 0s 5ms/step - accuracy: 0.6551 - loss: 0.6392 - val_accuracy: 0.6623 - val_loss: 0.6160
Epoch 16/200
20/20 — 0s 5ms/step - accuracy: 0.6803 - loss: 0.6285 - val_accuracy: 0.6688 - val_loss: 0.6133
Epoch 17/200
20/20 — 0s 5ms/step - accuracy: 0.6593 - loss: 0.6579 - val_accuracy: 0.6623 - val_loss: 0.6119
Epoch 18/200
20/20 — 0s 5ms/step - accuracy: 0.6569 - loss: 0.6207 - val_accuracy: 0.6558 - val_loss: 0.6099
Epoch 19/200
...
Epoch 199/200
20/20 — 0s 5ms/step - accuracy: 0.6902 - loss: 0.5440 - val_accuracy: 0.7792 - val_loss: 0.4985
Epoch 200/200
20/20 — 0s 5ms/step - accuracy: 0.7171 - loss: 0.5256 - val_accuracy: 0.7792 - val_loss: 0.5015
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

<keras.src.callbacks.history.History at 0x246a23eb0b0>
```