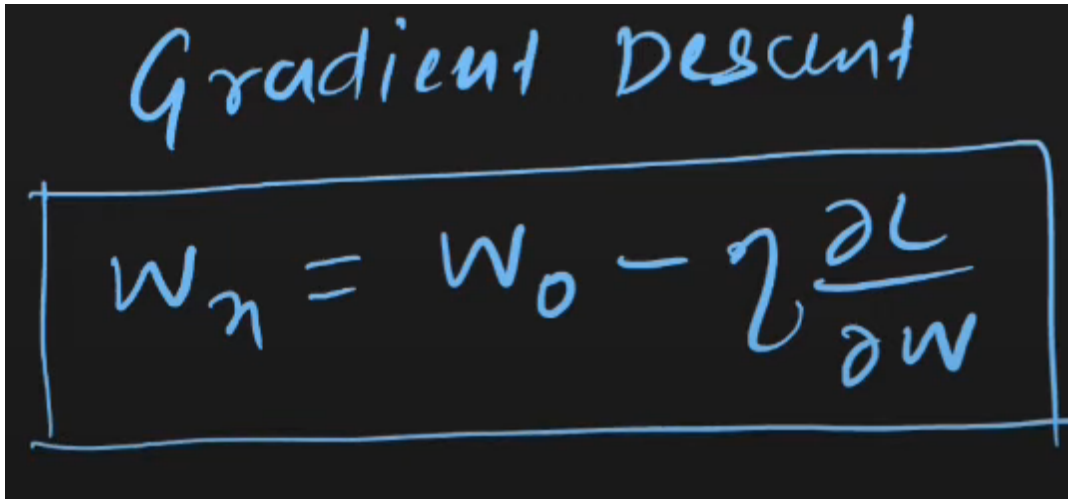


# Optimizers in Deep Learning


The optimizer we use in DL is **Gradient Descent**:



A handwritten equation on a blackboard background. The text 'Gradient Descent' is written at the top. Below it, a rectangular box contains the equation  $w_n = w_0 - \eta \frac{\partial L}{\partial w}$ .

## Challenges with GD:

- Deciding the value of learning rate ( $\eta$ )
- Cannot set a custom learning rate for individual weights and biases
- Local minima problem
  - We can get stuck in Local Minima
- Saddle point
  - There's a point where slope does not change
  - Therefore, weights don't get updated.



The following optimizers solve the above issues:

1. Momentum
2. Adagrad (Adaptive Gradient Algorithm)
3. NAG
4. RMSprop (Root Mean Square Propagation)
5. Adam

Key concept used in above optimizers → Exponentially weighted moving average

## Exponentially weighted moving average

- The **Exponentially Weighted Moving Average (EWMA)** is a smoothing technique that gives more weight to recent data while gradually "forgetting" older observations.

- It's a **weighted average**, but the weights **decay exponentially**
- EWMA gives more importance to recent weights
- As time passes, the weight of a point decreases exponentially


#### It's widely used in:

- **Time series forecasting** (e.g., stock prices, weather data).
- **Optimization algorithms** (e.g., Momentum in SGD, Adam, RMSprop).
- **Signal processing** (noise reduction).

 **Goal: Capture short-term trends while still considering past data.**

#### Why Use EWMA?

Problem	EWMA Fix
Too much noise in time-series	✅ Smooths fluctuations
Need quick reaction to new data	✅ Gives recent data more weight
Need memory of older data	✅ Never fully discards old values



Formula:

$$v_t = \beta \cdot v_{t-1} + (1 - \beta) \cdot \theta_t$$

- $\theta_t$ : Current observation (e.g., gradient, stock price).
- $\beta$ : Decay rate (typically **0.9**, **0.99**, etc.).
  - Higher  $\beta$  = smoother but slower to adapt.
  - Lower  $\beta$  = noisier but more responsive.

## Intuition

- Think of it as a **leaky average**, where older data decays exponentially.
- **Example:** If  $\beta=0.9$ , the weight of past observations decays like:

**Current=10%,**  
**1 step back=9%,**  
**2 steps back=8.1%,etc.**

## **β Value:**

Close to 1 (e.g., 0.9)	Fast response, recent data dominates ( <b>Most common</b> )
Close to 0 (e.g., 0.1)	Slow response, more smoothing

- **β closer to 1:** The weight on the most recent observation  $x_i$  is larger, meaning the EWMA is more sensitive to recent changes and reacts more quickly to new data.
  - This is typically used when you want to **prioritize recent values** more heavily.
- **β closer to 0:** The weight on the most recent observation is smaller, and the EWMA places more importance on older values, making it **less sensitive to recent changes**.
  - This is used when you want a **smoother** time series that reacts more slowly to new data.

## **EWMA in Python:**

```
df['Close'].ewm(alpha=0.9).mean()
```

```


0      9.710000
1      9.791818
2      9.934324
3      9.849424
4      9.854442
...
246    10.573514
247    10.435351
248    10.475535
249    10.695554
250    11.329555
Name: Close, Length: 251, dtype: float64

```

```

plt.plot(df['Date'], df['Close'])
df['Close'].ewm(alpha=0.3).mean().plot(color='r')

```




## OPTIMIZERS 📌

### 1. SGD with *Momentum* (🚀Speed)



**SGD with Momentum** is an extension of the basic **Stochastic Gradient Descent (SGD)** optimization algorithm.

- **Stochastic Gradient Descent (SGD) with Momentum** is like rolling a ball downhill—it uses past gradients to **accelerate convergence** and **escape local minima**.



## Key Benefits

- ✓ **Faster convergence:** Especially in flat or noisy loss landscapes.
- ✓ **Escapes local minima:** Momentum carries it through small bumps.
- ✓ **Less oscillation:** Smoother updates than vanilla SGD.
- ✓ **Solves High Curvature problem**

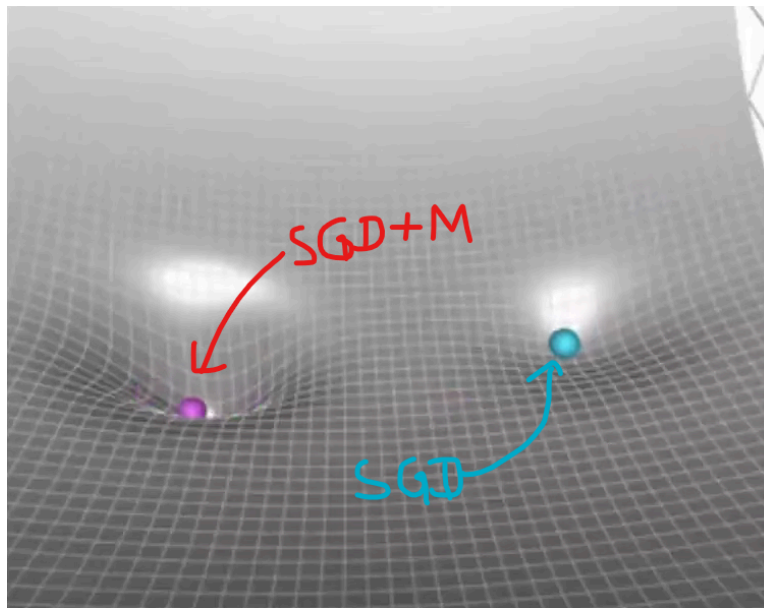
## ✗ Problem with vanilla SGD:

- Slow in narrow valleys

- Oscillates up and down in steep areas
- Gets stuck in local minima

### ✓ Momentum Fix:

- Adds **inertia** to the updates
- Helps SGD roll down like a **ball** with memory of past direction



### Formula:

$$v_t = \beta v_{t-1} + (1 - \beta) \nabla_{\theta} J(\theta_t)$$


$$\theta_{t+1} = \theta_t - \alpha v_t$$

- $v_t$ : Velocity (exponentially weighted average of past gradients).
- $\beta$ : Momentum hyperparameter (**0.9** is typical).
- $\alpha$ : Learning rate.

- 🙌 You use history of past velocities.



- It provides momentum



### $\beta$ Value:

- $\beta = 0 \rightarrow$  SGD
- $\beta = 1 \rightarrow$  Dynamic equilibrium. No decay
- $\beta = 0.9$  (Most common)

### Intuition (Ball in a Valley Analogy)

- Plain SGD: Like dropping a ball in a bumpy valley; it jumps side to side.
- SGD with Momentum: Ball gains **velocity** and rolls faster, **smoother**, and more directly toward the bottom.

### Python code:

```
from keras.optimizers import SGD

# SGD with momentum = 0.9
optimizer = SGD(learning_rate=0.01, momentum=0.9)

model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=
['accuracy'])
```

`learning_rate=0.01` → By default

**momentum= 0.0 (Default)**



## SGD vs SGD + Momentum

Feature	SGD	SGD + Momentum
Uses gradient only	✓	✓
Uses past gradients	✗	✓
Smooth updates	✗	✓
Converges faster	✗	✓
Handles ravines	✗	✓



## Why Momentum Beats Vanilla SGD

Scenario	Vanilla SGD	SGD + Momentum
Flat regions	Crawls slowly	Accelerates through
Noisy gradients	Oscillates wildly	Smoothens updates
Local minima	Gets stuck	Rolls past them

## Disadvantage of SGD + Momentum

- **Oscillations:** In regions of the loss surface where the gradient points in opposite directions (e.g., in valleys), SGD might oscillate back and forth.

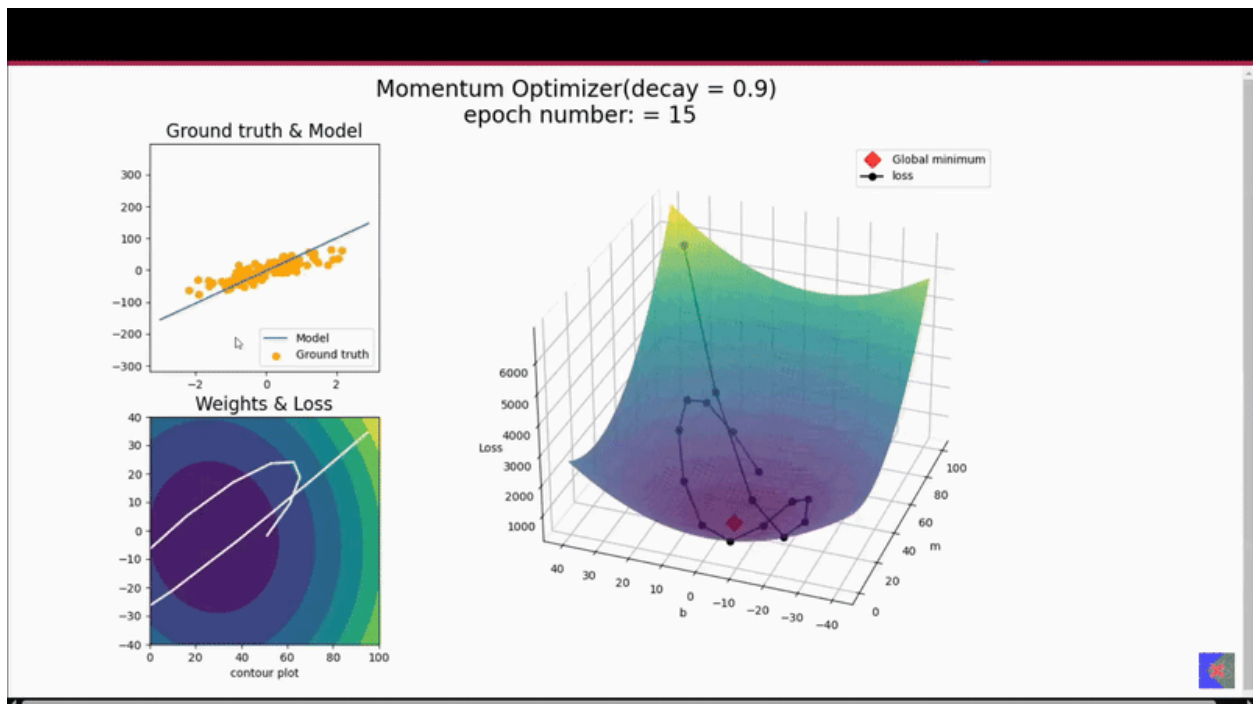
## 2. Nesterov Accelerated Gradient (NAG) 🐍

```
SGD(learning_rate=0.01, momentum=0.9, nesterov=True)
```

- Upgrade to SGD Momentum
- Mostly performs better than SGD Momentum

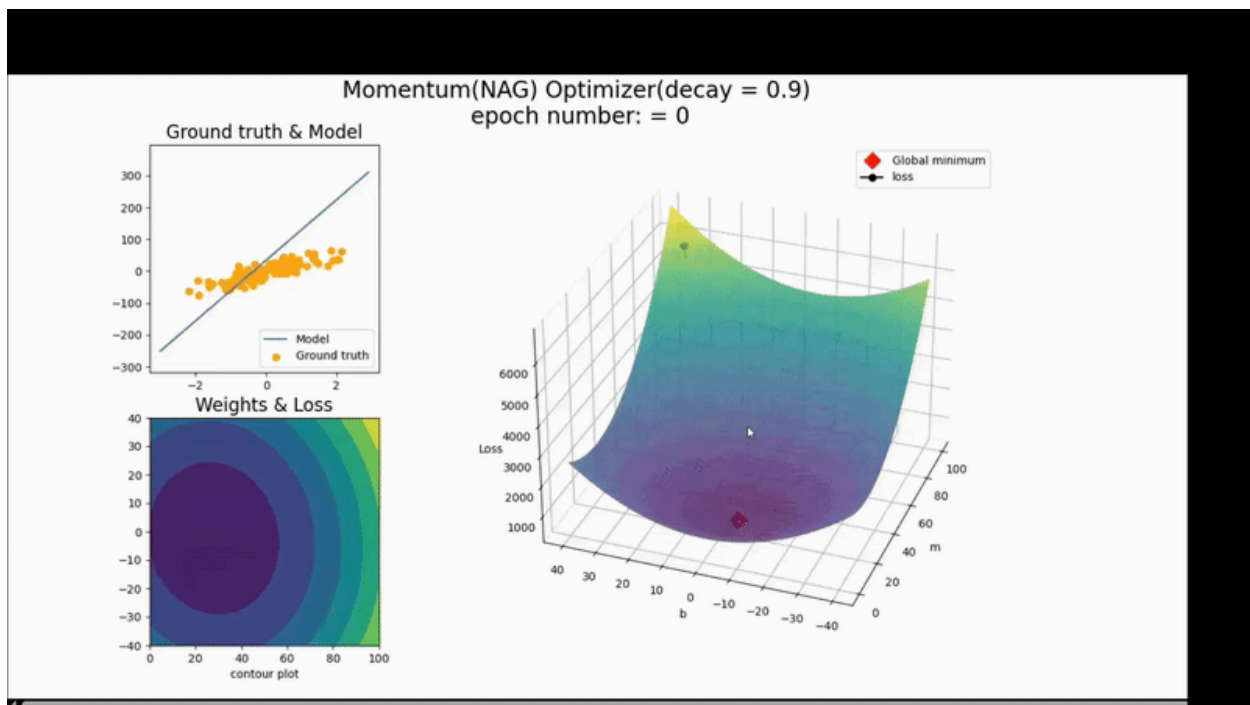
[attachment:e9887393-6286-49e9-8a30-a05a3d4d7f45:2025-04-05\\_23-46-22.mp4](#)

## SGD Momentum 🖐️🖐️



- The above 🖐️ graph is taking more epochs to reach to minima cuz the decay factor is 0.9.
  - i.e. it's giving more importance to the recent weights
- If we decrease the decay factor, the oscillations will reduce
- **NAG helps to decrease the oscillations by keeping the decay constant.**

## NAG with same decay i.e. 0.9 🖐️



**Goal: Make more accurate and faster updates by “looking ahead” before calculating the gradient.**

**“Before I take the next step, let me check the slope a little ahead of where I’m going.”**

## Why Use NAG?

Issue in SGD + Momentum	NAG Solution
Momentum blindly trusts past direction	NAG looks ahead before moving
Overshoots minima	NAG applies <b>correction</b> before stepping
Slower convergence	NAG converges <b>faster and smoother</b>

## NAG vs. Classic Momentum

Aspect	Classic Momentum	Nesterov Momentum
--------	------------------	-------------------

<b>Gradient Calculation</b>	At current position	At "lookahead" position
<b>Overshooting</b>	More likely	Reduced
<b>Convergence Speed</b>	Fast	<b>Faster</b>
<b>Stability</b>	Good	<b>Better</b>

## Update Rules – Step by Step

### ✓ Step 1:

- **Lookahead Step:** Instead of applying the momentum update directly on the current parameters

$\theta_t$ , NAG first performs a lookahead step to predict where the parameters will be after the update:

$$\tilde{\theta}_t = \theta_t - \gamma \cdot v_{t-1}$$

### ✓ Step 2: Compute Gradient at Lookahead

- Gradient is taken at the **lookahead position**, not current position.

$$\nabla J(\tilde{\theta}_t)$$

### ✓ Step 3: Update Velocity

$$v_t = \gamma \cdot v_{t-1} + \eta \cdot \nabla J(\tilde{\theta}_t)$$

### ✓ Step 4: Update Weights

$$\theta_{t+1} = \theta_t - v_t$$

## Python (Keras) Implementation

```
from keras.optimizers import SGD

# SGD with Nesterov Momentum
optimizer = SGD(learning_rate=0.01, momentum=0.9, nesterov=True)
```

### Explanation:

- `momentum=0.9` : Standard momentum value
- `nesterov=True` : Enables lookahead behavior

## Simple Analogy

Plain Momentum:

**"I'm going downhill fast. Keep going!"**

NAG:

**"I'm going downhill fast. But let me look a little ahead first to see if I should slow down."**

## When to Use NAG?

✓ Ideal for:

- Deep neural networks
- Highly non-convex loss surfaces
- Image classification, NLP, time series forecasting

## Disadvantage of NAG

- You can get stuck in local minima.
  - **Momentum** has enough power(momentum) to get out of the local minima.

## 🚫 When to Avoid NAG?

✗ When you're using **Adam** or **RMSProp**, which already handle momentum smartly.

✗ Very noisy gradients (Adam/RMSprop may be better).

✗ Extremely large batches (momentum matters less).

## 3. AdaGrad Optimizer

AdaGrad (Adaptive Gradient Algorithm) **adjusts the learning rate individually** for each parameter, based on how frequently it's been updated.



Adagrad is not used in Neural Networks.

🎯 Goal: Take larger steps for infrequent features, and smaller steps for frequent ones.

## Why Use AdaGrad?

### Problem with SGD:

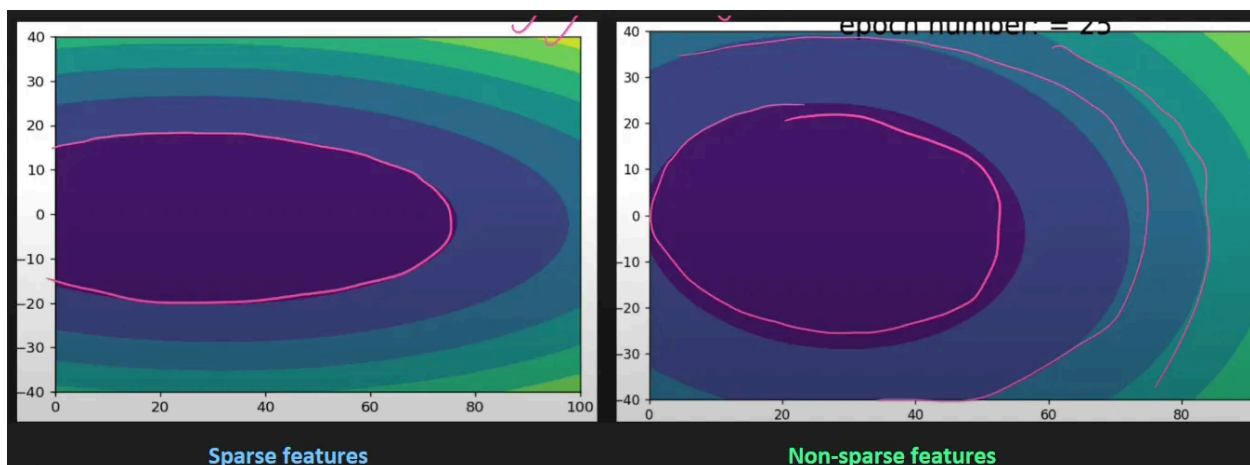

- Uses **same learning rate** for all weights → inefficient

### AdaGrad Fix:

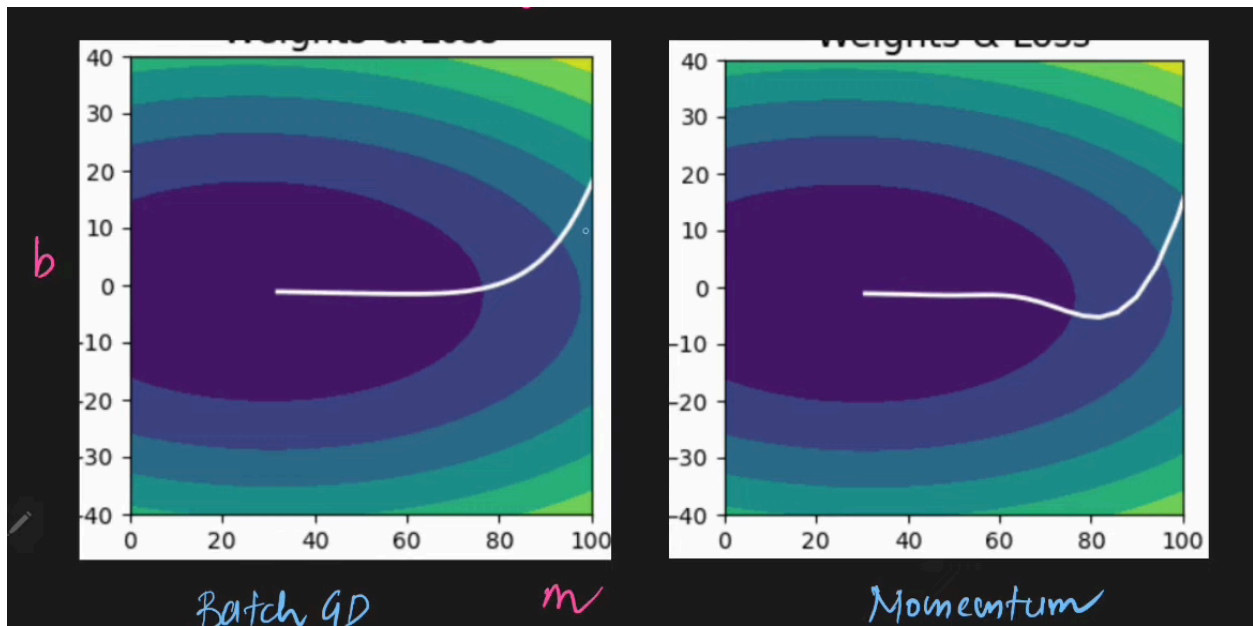
- Tracks the **sum of squared gradients** for each weight
- Adjusts learning rate **based on parameter history**

## ⚡ When to Use AdaGrad?

- When scale of input features is different
  - eg. CGPA (scale: 0 to 10) & salary (scale: in Lacs)
  - But we normalize the data in such case so Adagrad is not much useful.
- When features are sparse (most of the values are zero)







- The movement is slow for sparse columns cuz weight update is small
- To solve the above problem, adagrad uses adaptive LR 📌👉

📌 **Big gradient → Small Learning Rate (& vice versa)**

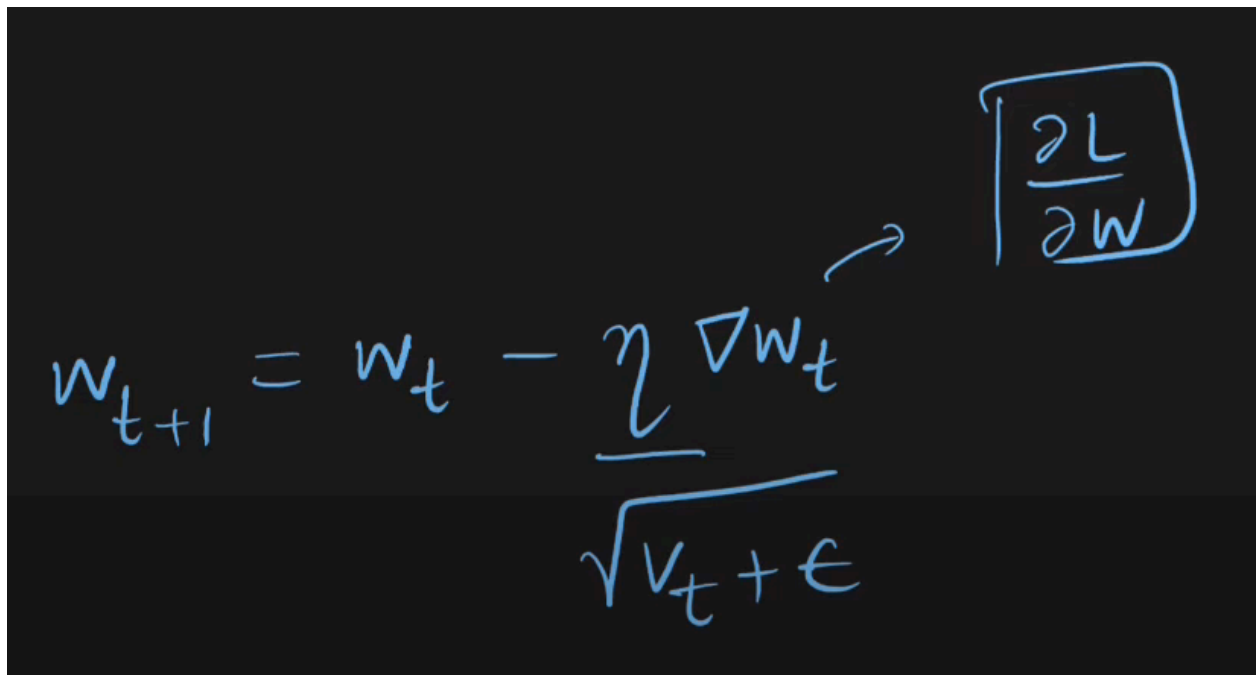
A handwritten formula on a blackboard: 
$$\underline{W} = W - \eta \frac{\partial L}{\partial W}$$
 The learning rate  $\eta$  is circled in red. A red arrow points from the text "Big gradient" to the denominator  $\frac{\partial L}{\partial W}$ , and another red arrow points from the text "Small Learning Rate" to the circled  $\eta$ .

**Formula:**

**Update Rule:**

$$\theta_{t+1} = \theta_t - \frac{\alpha}{\sqrt{G_t + \epsilon}} \cdot g_t$$

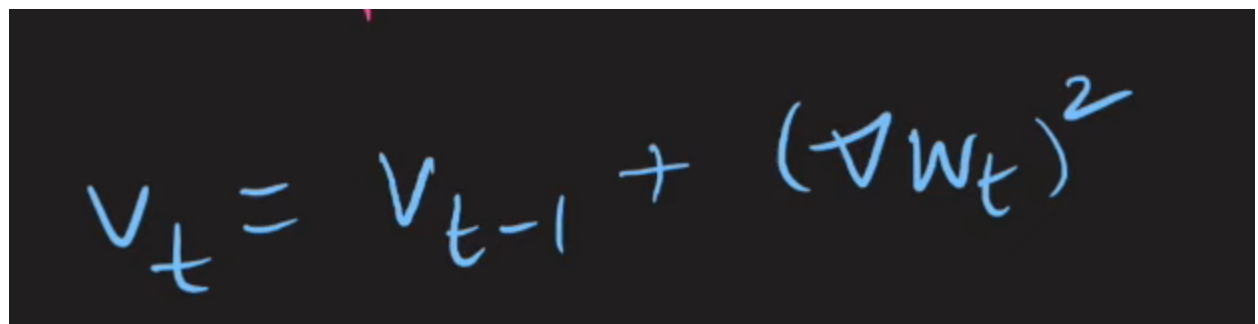
- $G_t$ : Sum of squared past gradients (per parameter).
- $\epsilon$ : Small constant ( $\sim 1e-8$ ) to avoid division by zero.



Handwritten update rule for weights  $w_t$  on a black background. The equation is  $w_{t+1} = w_t - \frac{\eta \nabla w_t}{\sqrt{v_t + \epsilon}}$ . An arrow points from the  $\nabla w_t$  term to a blue box containing the expression  $\frac{\partial L}{\partial w}$ .

$v_t$ : Sum of squared past gradients (per parameter)

$\epsilon$ : Small constant ( $\sim 1e-8$ ) to avoid division by zero.



Handwritten recursive formula for  $v_t$  on a black background:  $v_t = v_{t-1} + (\nabla w_t)^2$ .

## AdaGrad vs. SGD

Scenario	SGD	AdaGrad
<b>Sparse features</b>	Fails to train rare features	Excels
<b>Learning rates</b>	Fixed for all parameters	Adapts per parameter
<b>Convergence</b>	Slow for ill-conditioned data	Faster

## Disadvantages of AdaGrad

- AdaGrad can get close to solution but can never converge
  - **Reason:** we divide LR by  $v_t$ . It eventually gets large and LR decreases → Small Updates

## Behavior Summary:

Feature	Effect
Large gradient (frequent)	🚫 Step gets smaller
Small gradient (infrequent)	✅ Step stays large
Converges fast	✅ Initially
Long-term learning	❌ Learning rate may become too small to continue

## Python/Keras Implementation

```
from keras.optimizers import Adagrad

optimizer = Adagrad(learning_rate=0.01)

model.compile(optimizer=optimizer, loss='mse')
```

`learning_rate=0.001` (Default)

## 4. RMSProp (Root Mean Square Propagation)

- **Adaptive learning rate** optimization algorithm that **fixes AdaGrad's main weakness**: its learning rate decays too aggressively.
- **One of the best optimization techniques**



**RMSProp** was used before Adam.



**Goal:** Maintain stable, adaptive learning rates for each weight without vanishing updates.

- **Problem with AdaGrad:** The sum of squared gradients ( $G_t$ ) grows monotonically, causing learning rates to vanish over time.
- **RMSProp's Fix:** Replace  $G_t G_t$  with a **leaky average** (like momentum for gradients).



### Real-World Analogy

**AdaGrad:** "Every time I step, I get more cautious permanently."

**RMSProp:** "I remember the recent terrain and adjust my step size based on it—not the whole history."

### Update Rule:

$$E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta)g_t^2$$

$$\theta_{t+1} = \theta_t - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} \cdot g_t$$

- $E[g^2]_t$ : Exponentially weighted average of squared gradients.
- $\beta$ : Decay rate (typically **0.9**).
- $\epsilon$ : Small constant ( $\sim 1e-8$ ) for numerical stability.



### 3. RMSProp Update Rule – Step by Step

Let:

- $\theta$ : parameter (weight)
- $g_t$ : gradient at time t
- $E[g^2]_t$ : exponentially decaying average of squared gradients
- $\rho$ : decay factor (default = 0.9)
- $\eta$ : learning rate
- $\epsilon$ : small constant to avoid division by zero

### ✓ Step 1: Update Moving Average of Squared Gradients

$$E[g^2]_t = \rho \cdot E[g^2]_{t-1} + (1 - \rho) \cdot g_t^2$$

→ Like an Exponentially Weighted Moving Average (EWMA)

### ✓ Step 2: Parameter Update

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} \cdot g_t$$

## RMSProp vs. AdaGrad

Scenario	AdaGrad	RMSProp
Learning Rates	Vanish over time	Stabilize
Convergence	Stalls in deep nets	Works well
Hyperparameter	None	Decay rate ( $\beta$ )
Deep networks	✗ Not good	✓ Reliable

## Python/Keras Implementation

```
from keras.optimizers import RMSprop

optimizer = RMSprop(learning_rate=0.001, rho=0.9)

model.compile(optimizer=optimizer, loss='mse')
```

`learning_rate=0.001` (Default)

`rho=0.9` (Default)

## ⚡ Pro Tips

### 1. Default Hyperparameters:

- Learning rate ( $\alpha$ ): **0.001** (start small).
- Decay rate ( $\beta$ ): **0.9** (standard).

### 2. Combine with Momentum:

- RMSProp focuses on **gradient magnitude**; momentum handles **direction**.
- Result: **Adam optimizer** (next-gen default).

### 3. Use For:

- RNNs, non-convex problems.
- Replaced by Adam in most cases, but still useful for some tasks.

## 5. Adam (Adaptive Moment Estimation)

- Adam combines **RMSProp (adaptive learning rates)** and **Momentum (gradient history)** into one powerful optimizer.
- It's the **default choice** for most deep learning tasks due to its robustness and speed.

🎯 **Goal: Achieve fast convergence with stable updates in deep learning models.**

### 🧠 Why Use Adam?

Problem	Fix Adam Provides
Learning rate needs tuning	✅ Automatically adapts per weight
Momentum or RMSProp alone isn't enough	✅ Combines both
Noisy gradients in mini-batch training	✅ Smooth updates
Training unstable	✅ Stabilizes with bias correction

### 🔄 Update Rule – Step by Step

## Mathematical Formulation

04 August 2022 10:45

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{v_t} + \epsilon} * m_t$$

Bias correction epoch #

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\beta_1 = 0.9$$

$$\beta_2 = 0.99 \rightarrow \text{Keras}$$

where

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla w_t \rightarrow \text{momentum}$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla w_t)^2 \rightarrow \text{Adagrad}$$

- $\theta_t$ : current weight
- $g_t$ : current gradient
- $m_t$ : moving average of gradient (1st moment)
- $v_t$ : moving average of squared gradient (2nd moment)
- $\beta_1$ : decay for momentum (default = 0.9)
- $\beta_2$ : decay for RMS (default = 0.999)
- $\epsilon$ : small constant (default =  $1e-7$ )



✓ Step 1: Compute Gradients

$$g_t = \nabla_{\theta} J(\theta_t)$$

✓ Step 2: Update First Moment (Momentum)

$$m_t = \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$$

✓ Step 3: Update Second Moment (RMS)

$$v_t = \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$$

✓ Step 4: Bias Correction

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Corrects for initialization bias early in training.

### ✓ Step 5: Update Parameters

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \cdot \hat{m}_t$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad (\text{1st moment})$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad (\text{2nd moment})$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad (\text{bias correction})$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\theta_{t+1} = \theta_t - \frac{\alpha \cdot \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

### Adam vs. Other Optimizers

Optimizer	Adaptive LR?	Momentum?	Best For
SGD	✗	✗	Simple convex problems
SGD + Momentum	✗	✓	Deeper networks
AdaGrad	✓	✗	Sparse data
RMSProp	✓	✗	RNNs, non-convex landscapes
Adam	✓	✓	Most deep learning tasks

## Keras Code

```
from keras.optimizers import Adam

optimizer = Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999)

model.compile(optimizer=optimizer, loss='mse')
```

`learning_rate=0.001` (Default)

`beta_1=0.9` (Default)

`beta_2=0.999` (Default)

## Why Adam Dominates?

- ✓ **Adaptive Learning Rates:** Each parameter gets a custom step size.
- ✓ **Momentum Acceleration:** Smoothens noisy gradients.
- ✓ **Bias Correction:** Fixes early training instability.
- ✓ **Works Out-of-the-Box:** **Default hyperparameters** suit most problems.

## Use Cases

Task	Adam Use?
Deep neural networks	✓ Excellent
NLP & Transformers	✓ Standard choice
Image classification	✓ Common
RNNs & LSTMs	✓ Very effective
Tabular data	✓ Works well



## Real-World Analogy

Adam is like a self-driving car:

It adjusts its speed based on road slope (gradient), remembers recent direction (momentum), and avoids overreacting (bias correction).

## **When to Avoid Adam?**

- ✗ **Theoretical convex problems** (SGD may generalize better).
- ✗ **Extremely large batches** (momentum becomes less useful).

## **Summary**

- **Adam = RMSProp + Momentum + Bias Correction.**
- **Use For:** Almost all deep learning (CNNs, RNNs, Transformers).
- **Implementation:** 1 line in PyTorch/Keras.