Data Augmentation

- **Generates Data**

> Data augmentation artificially expands your training dataset by applying **random transformations** to images, improving model generalization and reducing overfitting.

🎯
Goal: Improve model generalization, reduce overfitting, and simulate real-world variations without collecting new data.

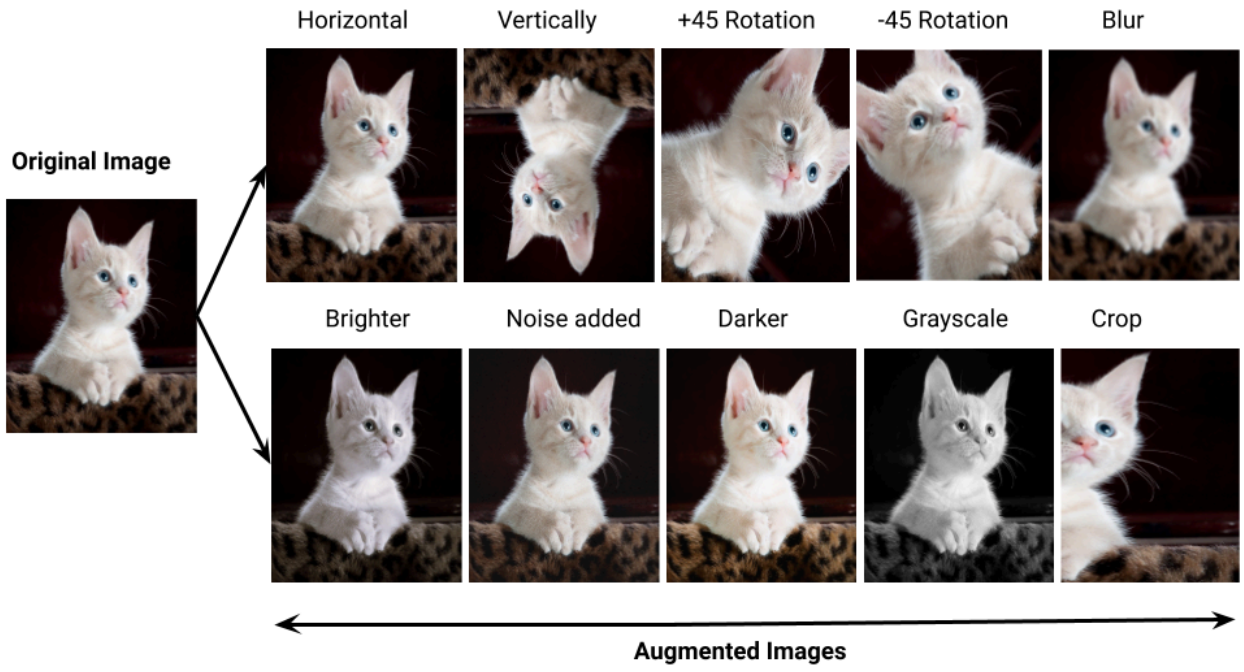

Why Use Data Augmentation?

Problem	How Augmentation Helps
Small dataset	🚀 Increases dataset size artificially

Problem	How Augmentation Helps
Overfitting	✅ Introduces variety, forces generalization
Lack of variation	✅ Simulates noise, rotation, etc.
Real-world robustness	✅ Prepares model for unseen cases

Common Augmentation Techniques (Image Data)

Technique	Keras Implementation	Use Case
Random Flip	`RandomFlip("horizontal")`	Mirror images (preserves semantics).
Random Rotation	`RandomRotation(0.2)`	Handles slight camera tilts.
Random Zoom	`RandomZoom(0.1)`	Simulates varying distances.
Random Contrast	`RandomContrast(0.2)`	Adjusts lighting conditions.
Random Crop	`RandomCrop(height, width)`	Focuses on partial views.

✅Best Practices

1. **Don’t Augment Validation/Test Data**:
2. **Use Moderate Values**
   - `rotation_range=20` (degrees), not 180° (cats don’t hang upside down!).
3. **Visualize Augmented Images**:
4. **Combine with Transfer Learning**:

🔁 Data Augmentation for Other Data Types

✅ Text (NLP):

Technique	Description
Synonym Replacement	Swap words with synonyms
Random Insertion	Insert words randomly
Random Deletion	Remove random words
Back Translation	Translate to another language and back

✅ Audio:

Technique	Description
Time Shift	Shift waveform left/right
Pitch Shift	Change audio pitch
Noise Injection	Add white noise
Speed Change	Slow down or speed up audio

🎯 When to Use?

Use Case	Use Data Augmentation?
Deep Learning (CV, NLP, Audio)	✅ Absolutely
Small datasets	✅ Mandatory
Test/Validation Set	❌ Never augment
Real-time inference	❌ No need

❌ Limitations & Warnings

- **Too much augmentation → garbage data → underfitting**
- Should **preserve label** (e.g., flipping digit 6 becomes 9 — bad)
- **Not all augmentations fit all problems** (e.g., rotation doesn't make sense for text)

Python Code:

```
from tensorflow.keras import layers

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.2),
    layers.RandomZoom(0.1),
    layers.RandomContrast(0.1),
```

```
])

# Usage in a model
model = tf.keras.Sequential([
    data_augmentation,  # Add augmentation as a layer
    layers.Conv2D(32, (3, 3), activation='relu'),
    # ... rest of the model
])
```

Advantages:

- Runs on GPU (much faster).
- Integrated directly into the model graph.