

Multilayer Perceptron (MLP)

```
from sklearn. neural_network import MLPClassifier
```

- An **MLP** is a **feedforward neural network**(type of **Artificial Neural Network (ANN)**) with **one or more hidden layers** between input and output.
- Unlike a single-layer perceptron, it can solve **non-linear problems** (like XOR) using **backpropagation** and **gradient descent**.

Structure of MLP

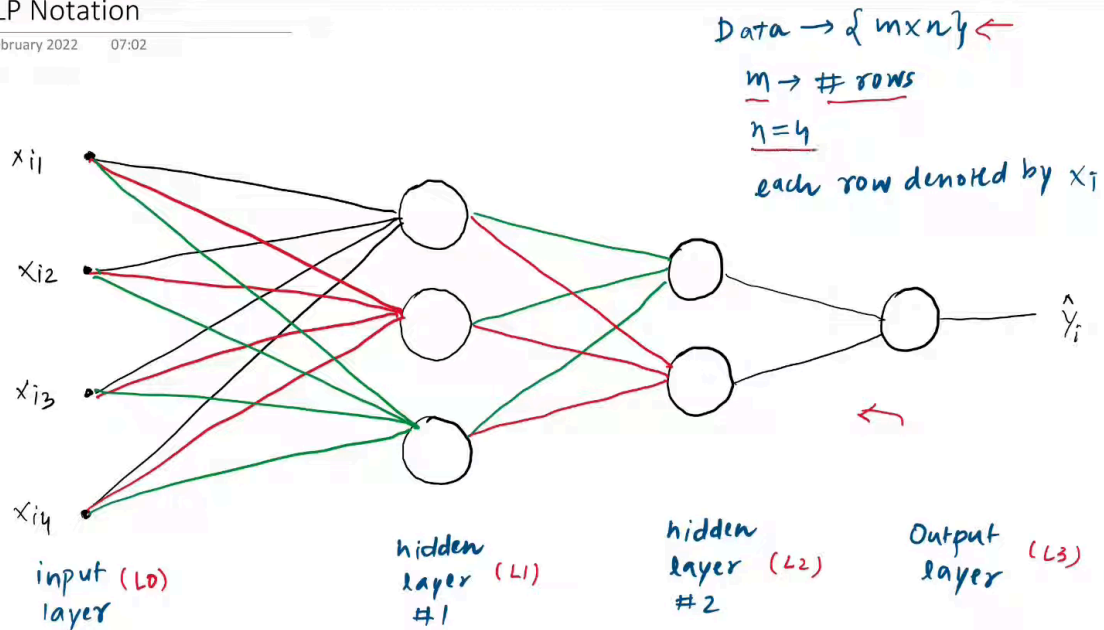
An MLP consists of three types of layers:

- **Input Layer**: Receives raw data (features).
- **Hidden Layer(s)**: Performs computations and extracts patterns.
- **Output Layer**: Produces the final prediction



MLP Notation

26 February 2022 07:02




- Find out how much trainable parameters you have
 - When you'll train this algorithm, how many **biases and weights** will be calculated
- In above image:
 - Weights:**
 - $3 \times 4 = 12$
 - $3 \times 2 = 6$
 - $2 \times 1 = 1$
 - 1
 - Biases:**
 - 3
 - 2
 - 1
- Add all the 🖐 above weights & biases:

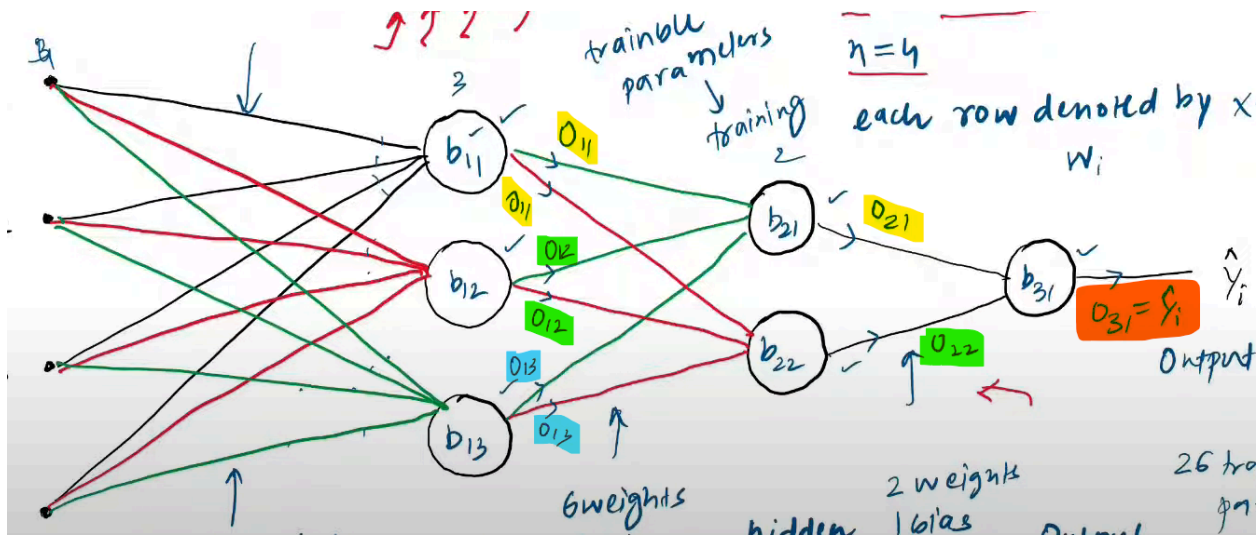
- $12+6+1+3+2+1=26$
- There are 26 trainable parameters.

MLP Notation

- **Bias** $\rightarrow b_{ij}$
 - i = layer number
 - j = node in the layer i



- 🖐 $b_{11} \rightarrow$ First layer, first node
- **Output** $\rightarrow o_{ij}$



• Weights

Notation: W_{ij}^h

where,

i = From which node weight is passing to the next layer's node.

j = To which node weight is arriving.

h = Layer in which weight is arriving.

Example:

W_{11}^1 = Weight passing into 1st node of 1st layer of 1st node of the previous layer.

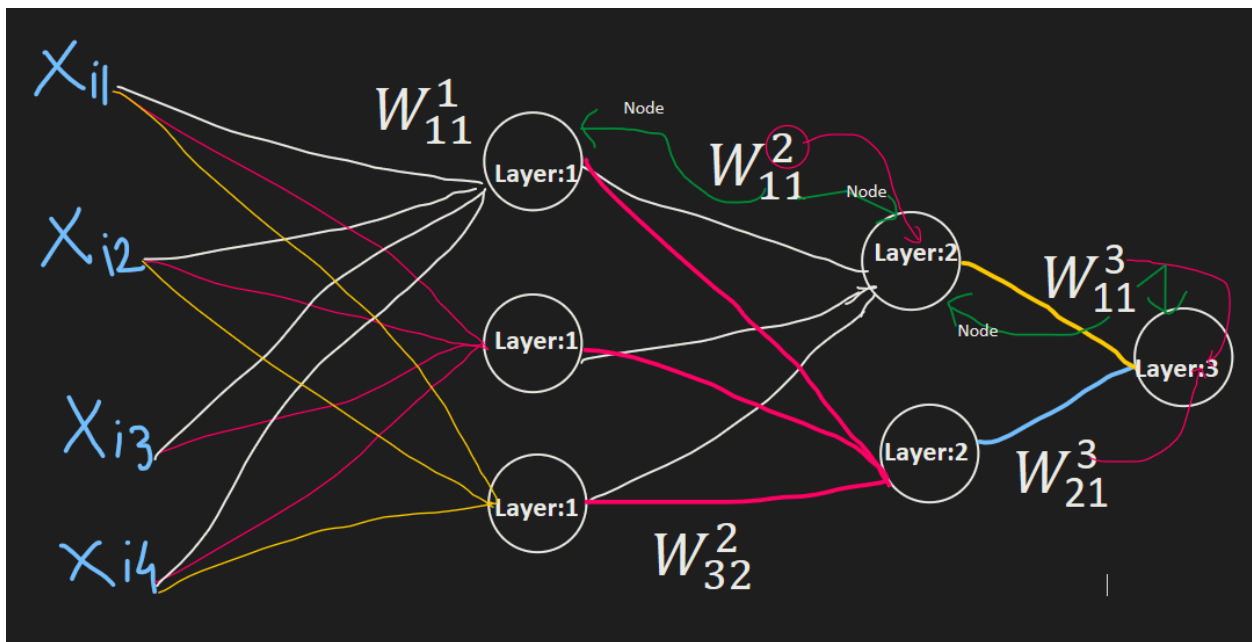
W_{23}^1 = Weight passing into the 3rd node of the 1st hidden layer from the 2nd node of the previous layer.

W_{45}^1 = Weight passing into the 5th node of the 2nd hidden layer from the 4th node of the previous layer.

3. Notations for Outputs:

$$W_{ij}^k$$

- **k**: **Layer** in which weight is arriving
- **i**: Current layer's **node**
- **j**: **Node** to which it's arriving



MLP vs. Single-Layer Perceptron

Feature	Single-Layer Perceptron	MLP
Layers	Input → Output	Input → Hidden → Output
Non-Linear?	Fails (XOR problem)	Solves non-linear problems
Training	Perceptron Trick	Backpropagation + Gradient Descent

Use Case	Linear classification	Complex tasks (images, NLP)
----------	-----------------------	-----------------------------


Loss Functions used:


Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}$ (for probabilities)

ReLU: $f(z) = \max(0, z)$ (for deep networks)

Softmax: Used in multi-class classification.

- In MLP, we **combine multiple perceptrons** & **smoothen the boundary**.





- We can manipulate importance of 2 perceptrons by multiplying them by weights

$$0.7 \times 10 + 0.8 + 5 = z_1$$

- Here, 10 & 5 are weights of 1st & 2nd perceptron
- After this, we send this to sigmoid function.


You can also add a bias.

$$[z] = 0.7 \times 10 + 0.8 \times 5 + 3$$

- 🙌 3 is bias.

Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}$ (for probabilities)

- If $\sigma \geq \text{threshold (0.5)} \rightarrow 1$
- If $\sigma < \text{threshold} \rightarrow 0$




You can add multiple nodes.

- If you add more nodes, more it'll help in defining the non-linear boundary

Add Multiple output nodes in output layer in case of **multi-class classification** problem.

You can add No. of hidden layers.

- You can make more complex decision boundary by adding more layers.



Backpropagation (Learning Process)

- Computes the error (difference between predicted and actual output).
 - Uses **Gradient Descent** to update weights and minimize error.
 - Adjusts weights using **chain rule of differentiation**.
1. **Forward pass:** Input data is passed through the network's layers to generate an output
 2. **Error calculation:** The difference between the input and output is calculated
 3. **Backward pass:** The loss is calculated backwards, layer by layer
 4. **Weights update:** The weights are updated based on the gradient calculated in the backward pass

Python code for MLP:

```

from sklearn.neural_network import MLPClassifier
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Generate dataset
X, y = make_moons(n_samples=1000, noise=0.2, random_state=42)

# Split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define MLP model
mlp = MLPClassifier(hidden_layer_sizes=(10, 5), max_iter=1000)

# Train model
mlp.fit(X_train, y_train)

# Predict
y_pred = mlp.predict(X_test)

# Accuracy
print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")

```

Accuracy: 0.89

hidden_layer_sizes (Defining the Network Architecture)

What it does?

- Specifies the **number of hidden layers** and **neurons per layer** in the MLP.
- Default: `(100,)` → **1 hidden layer with 100 neurons**

hidden_layer_sizes	Neural Network Architecture
(100,)	1 hidden layer with 100 neurons
(50, 30)	2 hidden layers (50 neurons, 30 neurons)
(10, 20, 30)	3 hidden layers (10, 20, 30 neurons)
(150, 100, 50, 25)	4 hidden layers (150 → 100 → 50 → 25 neurons)

- ◆ **More neurons = More capacity** (but may overfit).
- ◆ **More layers = More complexity** (but harder to train).

solver (Choosing the Optimization Algorithm)

What it does?

- Determines how the model updates weights during training (optimizer).
- Default: **'adam'**

Available Options & When to Use Them?

Solver	Algorithm Type	Best For
'adam' (Default)	Adaptive Moment Estimation (fast, adaptive learning rate)	Large datasets, complex problems
'sgd'	Stochastic Gradient Descent (simple, but slow)	Large-scale problems, online learning
'lbfgs'	Limited-memory BFGS (second-order optimization)	Small datasets, faster convergence

- **Use 'adam' (default)** if unsure – works well in most cases.
- **Use 'sgd'** when fine-tuning the learning rate manually.
- **Use 'lbfgs'** for smaller datasets with fast convergence.


Forward Propagation (Forward Pass)

- Forward propagation is how an **MLP makes predictions** by passing input data through its layers.

- It involves passing the data through the layers of the network, performing computations at each layer, and generating the predicted result.



Goal: To compute the predicted output \hat{y} given input X and weights W .



Forward Propagation Steps

1 Input Layer → First Hidden Layer

Input: X

- These are typically the raw data, like pixel values of an image or measurements of some object.

Weighted Sum (Linear Transformation):

- For each neuron in the hidden layer (and subsequent layers), the input is multiplied by a set of weights. Each connection between neurons has an associated weight

The weighted sum for each neuron is calculated as:

$$z = \sum (x_i \cdot w_i) + b$$

- x_i represents the input values (features),
- w_i represents the weights,
- b is the bias term.

Activation (e.g., ReLU, Sigmoid, or TanhU):

$$a^1 = \text{ReLU}(z^1)$$

Sigmoid activation:

$$a = \sigma(z) = \frac{1}{1 + e^{-z}}$$

where σ is the sigmoid function, and z is the weighted sum.

2 Hidden Layer → Output Layer

- The same process repeats for each hidden layer.
- The final layer computes the output using an activation function (e.g., **Sigmoid for classification, Softmax for multi-class**).

3. Example of Forward Propagation (Simple 2-Layer Network)

Assume a neural network with:

- **Input Layer:** 2 neurons
- **Hidden Layer:** 2 neurons
- **Output Layer:** 1 neuron

Given Inputs:

$$X = [x_1, x_2]$$

Weight Matrices:

$$W_1 = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix}, W_2 = [w_{31} \quad w_{32}]$$

Bias Vectors:

$$b_1 = \begin{bmatrix} b_1 \\ b_2 \end{bmatrix}, b_2 = b_3$$

Step 1: Compute Hidden Layer Activations

$$Z_1 = W_1 \cdot X + b_1$$

$$A_1 = f(Z_1) \quad (\text{ReLU, Sigmoid, etc.})$$

Step 2: Compute Output Layer Activation

$$Z_2 = W_2 \cdot A_1 + b_2$$

$$\hat{y} = f(Z_2) \quad (\text{Sigmoid, Softmax, etc.})$$