

# Deep RNNs (Stacked RNNs)





## What is a Deep RNN?

A **Deep RNN** is simply a **stack** (or **multi-layer**) of RNN layers.

In a **basic RNN**, you have **one recurrent layer**, which processes sequences step by step. But with **deep RNNs**, you stack **multiple RNN layers** on top of each other, allowing the model to learn **more complex patterns** and **abstract representations**.

Input ➡ RNN ➡ RNN ➡ RNN ➡ Output

## Why use Deep RNNs?

| Advantage  | Meaning                                    |
|--|--|
|  More representation power      | Learns deeper patterns in sequences        |
|  Better temporal modeling       | Captures long-term dependencies better     |
|  Handles complex language tasks | Like translation, speech recognition, etc. |
|  Captures the hierarchy         |  |

## Deep LSTM in Keras

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense

model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=64))
model.add(LSTM(128, return_sequences=True)) # 1st LSTM layer
model.add(LSTM(64)) # 2nd LSTM layer
model.add(Dense(1, activation='sigmoid')) # Output layer

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accr
```

```
acy'])  
model.summary()
```

`input_dim` → How many unique words we have

`output_dim` → Size of each word vector



**Note:** `return_sequences=True` is required for all layers except the last, so the output is still a sequence that the next RNN layer can read.

## Sentiment analysis:

```
import tensorflow as tf  
from tensorflow.keras.datasets import imdb  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense, LSTM, GRU
```

```
# Load the IMDB dataset  
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=10000)  
  
# Pad sequences to have the same length  
x_train = pad_sequences(x_train, maxlen=100)  
x_test = pad_sequences(x_test, maxlen=100)
```

## Build a model:

```
# Define the LSTM model  
model = Sequential([  
    Embedding(10000, 32),  
    LSTM(5, return_sequences=True),
```

```
LSTM(5),
Dense(1, activation='sigmoid')
])

model.summary()
```

## Compile:

```
# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

## Fit:

```
# Train the model
history = model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.2)
```

```
Epoch 1/5
625/625 — 19s 27ms/step - accuracy: 0.6614 - loss: 0.5871 - val_accuracy: 0.8258 - val_loss: 0.3948
Epoch 2/5
625/625 — 16s 26ms/step - accuracy: 0.8882 - loss: 0.2965 - val_accuracy: 0.8396 - val_loss: 0.3659
Epoch 3/5
625/625 — 16s 25ms/step - accuracy: 0.9260 - loss: 0.2044 - val_accuracy: 0.8402 - val_loss: 0.3792
Epoch 4/5
625/625 — 16s 25ms/step - accuracy: 0.9507 - loss: 0.1505 - val_accuracy: 0.8372 - val_loss: 0.4409
Epoch 5/5
625/625 — 16s 25ms/step - accuracy: 0.9638 - loss: 0.1155 - val_accuracy: 0.8298 - val_loss: 0.4808
```

## You can also use GRU:

```
# Define the GRU model
model = Sequential([
    Embedding(10000, 32),
    GRU(5, return_sequences=True),
    GRU(5),
    Dense(1, activation='sigmoid')
```

```
])
```

```
model.summary()
```

```
# Compile the model
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Train the model
```

```
history = model.fit(x_train, y_train, epochs=5, batch_size=32, validation_split=0.2)
```

```
Epoch 1/5
625/625 ————— 19s 27ms/step - accuracy: 0.6448 - loss: 0.6085 - val_accuracy: 0.8254 - val_loss: 0.4096
Epoch 2/5
625/625 ————— 16s 26ms/step - accuracy: 0.8750 - loss: 0.3155 - val_accuracy: 0.8382 - val_loss: 0.3747
Epoch 3/5
625/625 ————— 16s 26ms/step - accuracy: 0.9180 - loss: 0.2229 - val_accuracy: 0.8380 - val_loss: 0.3786
Epoch 4/5
625/625 ————— 16s 25ms/step - accuracy: 0.9479 - loss: 0.1555 - val_accuracy: 0.8334 - val_loss: 0.4212
Epoch 5/5
625/625 ————— 17s 26ms/step - accuracy: 0.9653 - loss: 0.1119 - val_accuracy: 0.8350 - val_loss: 0.4947
```

## When to use Deep RNNs?

- Complex tasks
  - eg. Speech Recognition, Machine Translation
- Large datasets
- Sufficient Computational power
- Not satisfied with simple models