

LSTM (Long Short Term Memory)

◆ Summary (Quick Orientation)

- LSTM is a special kind of **Recurrent Neural Network (RNN)**.
- It's designed to **remember long-term dependencies** and **solve the vanishing gradient problem**.
- It does this using **gates** that control what to **keep**, **forget**, and **output**.


Think of it as a memory cell with smart gates that decide:

- **What information to keep in memory**
- **What to throw away**
- **What to send to the output**

Invented by Hochreiter & Schmidhuber in 1997.

Here, you maintain 2 Paths:

1. Short term
 2. Long term
- **If you find a thing important, you put it in long term path.**
 - If you don't remove it, it will be there in the long term path till the end.

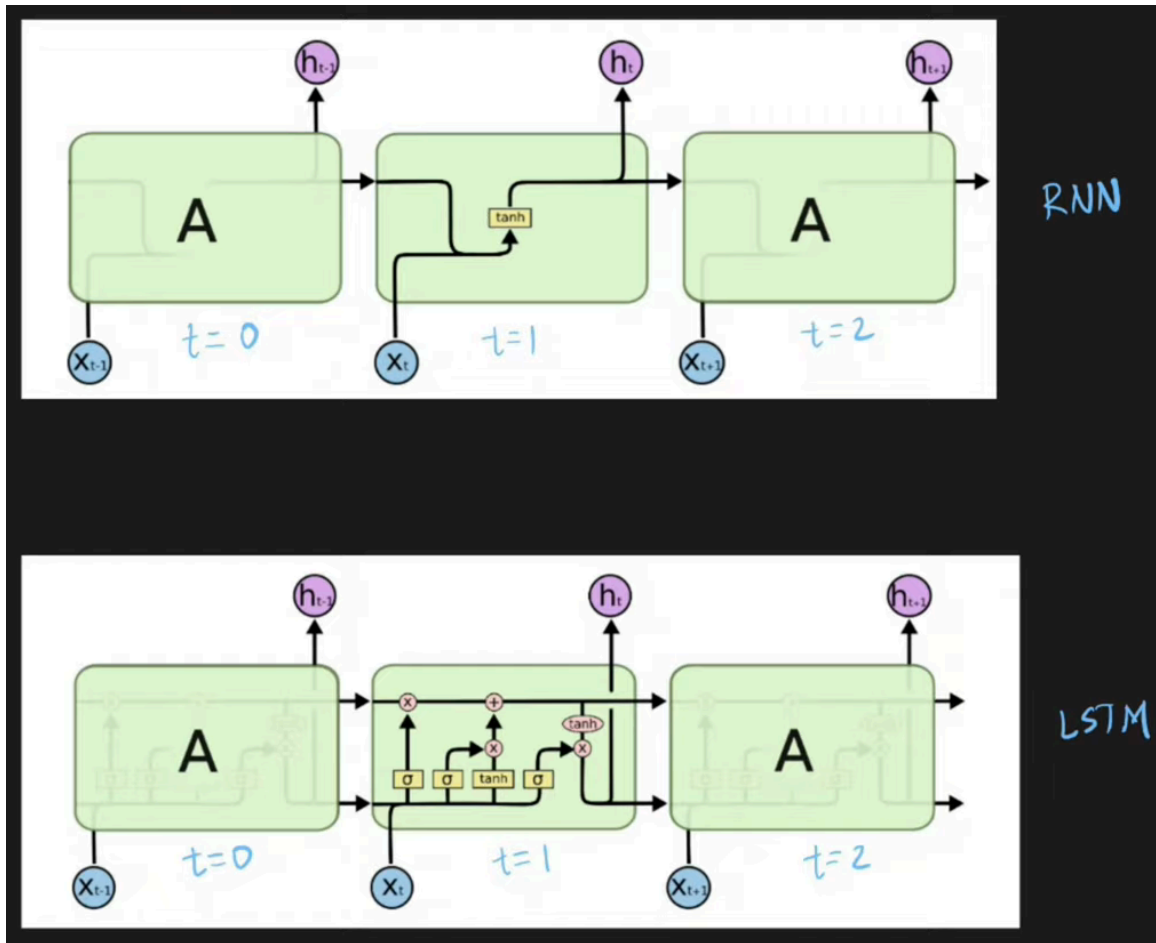


RNN has 2 paths.

Communication between these 2 happen with the architect.



LSTM Cell Structure



Each LSTM unit has:

1. **Cell state** C_t :

- **Long-term memory** (like a conveyor belt of info).
- Main memory that flows through time

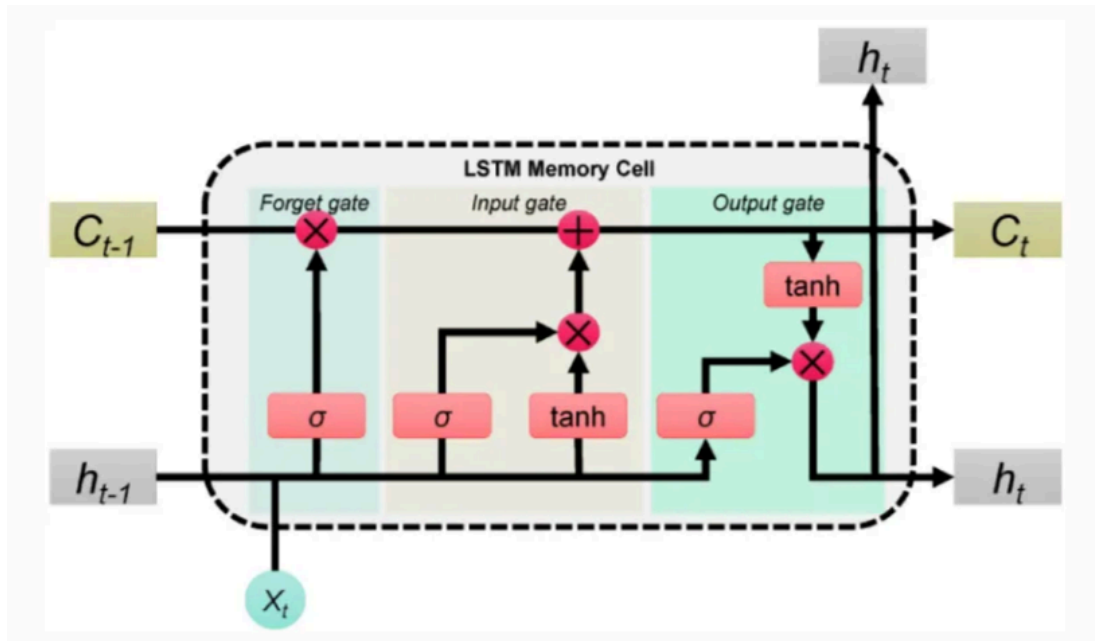
2. **Hidden state** h_t :

- **Short-term memory** (output at each step).
- Output passed to next time step

3. **Three gates:**

- Forget Gate:** Decides **what information to discard** from the cell state

- b. **Input Gate:** Decides **what new information to store** in the cell state
- c. **Output Gate:** Decides **what to output** based on the cell state



- **Input:**
 - $C_{t-1} \rightarrow$ Long Term Memory of previous Time step (Cell State)
 - $h_{t-1} \rightarrow$ Short Term Memory of previous Time step
 - $X_t \rightarrow$ Current word
- **Output:**
 - Long Term Memory for next step
 - Short Term Memory for next step
- **Inside:**
 - Update Long Term Memory
 - Create Short Term Memory

◆ 1. Forget Gate (f_t)

- **Purpose:** Decide what **past memory** to forget.

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Uses a **sigmoid (σ)** function \rightarrow output between 0 and 1
- If $f_t = 0 \rightarrow$ forget that info completely
- If $f_t = 1 \rightarrow$ keep it fully

◆ 2. Input Gate (i_t) + Candidate Memory (\tilde{C}_t)

Purpose: Decide what **new information** to add to memory.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \tilde{C}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

- i_t : input gate (what to add)
- \tilde{C}_t : candidate content (new data to be stored)

◆ 3. Update Cell State (C_t)

Purpose: Update the main memory cell.

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$


- Old memory is **partially forgotten**
- New memory is **partially added**
- This controlled flow keeps memory stable over time

◆ 4. Output Gate (o_t)

Purpose: Decide what to output at this time step.

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)h_t = o_t * \tanh(C_t)$$

- Final hidden state h_t is what's **sent to the next step**
- Also used as the **output** if needed



- The dimensions of above vectors are same.

🧠 Summary Flow of LSTM Cell at Time t :

1. Decide what to forget → f_t
2. Decide what to remember → i_t, \bar{C}_t
3. Update memory → C_t
4. Decide what to output → h_t

Applications of LSTMs

1. Sequence prediction
2. Machine translation
3. Speech recognition
4. Time series forecasting
5. Text generation
6. Video analysis

Hyperparameter Considerations

1. **Number of units:** Typically between 64-1024
2. **Number of layers:** Usually 1-4
3. **Dropout:** Recurrent dropout is often helpful
4. **Initialization:** Orthogonal initialization for recurrent weights

Python Code for LSTM:

Import libraries:

```
# Import necessary tools
from tensorflow.keras.models import Sequential # Basic neural network container
from tensorflow.keras.layers import Embedding, LSTM, Dense # Layers we'll use
from tensorflow.keras.preprocessing.text import Tokenizer # For text handling
from tensorflow.keras.utils import pad_sequences # To make sequences same length
import numpy as np
```

Data:

```
texts = ["I loved this movie", "Hated the film", "Best movie ever", "Worst experience"]  
labels = np.array([1, 0, 1, 0]) # 1=positive, 0=negative
```

Step 1: Prepare the text data:

```
# Step 1: Prepare the text data  
tokenizer = Tokenizer(num_words=10000) # Keep top 10,000 words  
tokenizer.fit_on_texts(texts) # Learn all words in our texts  
sequences = tokenizer.texts_to_sequences(texts) # Convert words to numbers  
data = pad_sequences(sequences, maxlen=100) # Make all reviews 100 words long
```

Step 2: Build the model

```
# Step 2: Build the model  
model = Sequential() # Create empty model  
  
# Add layers one by one:  
# 1. Embedding: Turns word numbers into meaningful vectors  
model.add(Embedding(input_dim=10000, # How many unique words we have  
                    output_dim=128)) # Size of each word vector  
  
# 2. LSTM layer: Understands sequences in the text  
model.add(LSTM(units=64)) # 64 memory units  
  
# 3. Dense layer: Final decision maker (positive/negative)  
model.add(Dense(1, activation='sigmoid')) # 1 output: 0-1 probability
```

64 → The number of neurons (or memory cells) inside the LSTM layer.

Compile:

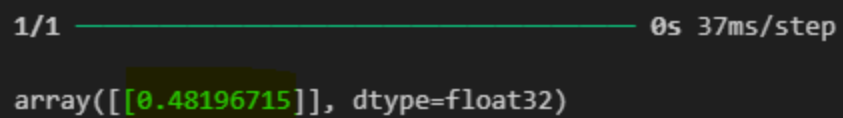

```
model.compile(optimizer='adam',    # Smart learning algorithm
              loss='binary_crossentropy', # How to measure errors
              metrics=['accuracy']) # Track correct guesses
```

Fit:

```
# Step 4: Train the model
model.fit(data, labels, # Our prepared data and answers
          epochs=5,    # How many times to see all data
          batch_size=32) # Process 32 reviews at once
```

Predict:

```
# Now the model can predict sentiments!
test_text = ["This film was okay"]
test_seq = tokenizer.texts_to_sequences(test_text)
test_data = pad_sequences(test_seq, maxlen=100)
prediction = model.predict(test_data) # Returns something like [[0.65]] (65%
positive)
prediction
```

A terminal window with a dark background. At the top, it shows '1/1' followed by a green progress bar and '0s 37ms/step'. Below this, it displays the output 'array([[0.48196715]], dtype=float32)' where the value '0.48196715' is highlighted in green.

```
1/1 ————— 0s 37ms/step
array([[0.48196715]], dtype=float32)
```